

# Log-concave Sampling (Part 1)

Marios Papachristou

GeomScale, NTUA

June 1, 2020



GeomScale



Mentors: A. Chalkis (NKUA), V. Fisikopoulos (NKUA), E. Tsigaridas (INRIA)

Homepage: [https://github.com/GeomScale/volume\\_approximation](https://github.com/GeomScale/volume_approximation)

# About today's talk / tutorial

Today's talk will concentrate on

## **Sampling from high-dimensional log-concave densities**

- 1 Introduction to log-concave sampling.
- 2 ODE Solvers.
- 3 Boundary Oracles.

# Google Summer of Code 2020

The current GSoC project aims to provide implementations (and theoretical insights) to log-concave sampling problems for the GeomScale project.

## Milestones

- ① Milestone I (ODE Solvers)
  - Implement ODE solvers (Euler, Runge-Kutta, Collocation, etc.)
  - Efficiently address boundary oracles
- ② Milestone II (Samplers)
  - Implement samplers (HMC, Langevin etc.).
  - Provide theoretical guarantees on truncated settings.
- ③ Milestone III (R bindings)
  - Port C++ functionality of

Today's talk will mostly concentrate on **Milestone I**.

Our project involves taking samples from distributions with probability density functions of the form

$$\pi(x) \propto \exp(-f(x)) \quad x \in K$$

where  $K$  is either: (a)  $\mathbb{R}^d$ , or (b) a convex body, and  $f$  is a convex function that is  $L$ -smooth and  $m$ -strongly convex.

# Convex Functions I

A domain  $K$  is convex iff (if and only if) for all  $x, y \in K$  it holds that for all  $t \in [0, 1]$

$$tx + (1 - t)y \in K$$

The domain  $K$  is a convex body iff it is convex, closed and bounded.  
A function  $f : K \rightarrow \mathbb{R}$  is convex iff for all  $x, y \in K$  we have that for all  $t \in [0, 1]$

$$f(tx + (1 - t)y) \leq tf(x) + (1 - t)f(y)$$

Convex functions have some very nice properties, and their use is widespread in optimization.

## Convex Functions II

If the function is twice differentiable with gradient  $\nabla f$  and Hessian matrix  $\nabla^2 f$  then

- We say that  $f$  is  $L$ -smooth iff  $\|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|$  or  $\nabla^2 f(x) \preceq L \cdot I_d$ .
- We say that  $f$  is  $m$ -strongly convex iff

$$f(y) \geq f(x) + \langle \nabla f(x), y - x \rangle + \frac{m}{2} \|x - y\|^2$$

or  $\nabla^2 f(x) \succeq m \cdot I_d$ .

- We define the **condition number** of  $f$  to be the ratio of max/min eigenvalues of the Hessian, that is  $\kappa = L/m$ .

# Log-concave Sampling I

Our goal is sampling from  $\pi(x) \propto \exp(-f(x))$ .

Directly sampling from  $\pi(x)$  is very difficult since one has to account for the normalization constant  $\int_K \exp(-f(x)) dx$  which is in general **intractable**.

**Idea:** The distribution  $\pi(x)$  can be thought as the stationary measure of a Markov Chain that is  $\pi(x) = \lim_{k \rightarrow \infty} \pi_k(x)$ .

# Log-concave Sampling II

One of the first algorithms to do it is the Metropolis-Hastings Algorithm. The general idea of Metropolis Hastings is

- 1 Assume that you are at a state  $x$
- 2 Perform a transition to a new nearby state  $y$  and make a proposal
- 3 Accept the proposal to move to  $y$  with probability (Metropolis Filter)

$$\min \left\{ 1, \frac{a(x, y)\pi(y)}{a(y, x)\pi(x)} \right\}$$

where  $a$  is a transition probability function.

It can be shown analytically that the above process converges to a stationary distribution  $\pi(x)$ .

**Intuition when  $a(x, y) = a(y, x)$ :** The sampler has incentive to move towards higher-density areas (but lower density areas are also allowed)



# Sampling in a continuous setting

# Hamiltonian Monte Carlo I

The state-space is continuous and the samples can be proposed via solving Hamilton's equations for a particle with position  $x$  and velocity  $v$  under a conservative potential  $f(x)$  that applies a force  $-\nabla f(x)$ .

The Hamiltonian of the particle is defined as

$$H(x, v) = \frac{1}{2} \|v\|^2 + f(x)$$

and Hamilton's equations simulate the particle's behaviour

$$\dot{x} = v$$

$$\dot{v} = -\nabla f(x)$$

# Hamiltonian Monte Carlo II

We start by choosing a direction  $v \sim \mathcal{N}(0, I_d)$  and simulate one/many steps of the ODE arriving at a proposal  $(\tilde{x}, \tilde{v})$ .

The Metropolis Filter in this case for a proposal  $(\tilde{x}, \tilde{v})$  given a state  $(x, v)$  is  $\min\{1, \exp(H(\tilde{x}, \tilde{v}) - H(x, v))\}$ .

Ideally note that  $\dot{H} = \langle \nabla_{x,v} H, (\dot{x}, \dot{v}) \rangle = \langle (v, \nabla f(x)), (-\nabla f(x), v) \rangle = 0$  and hence the Metropolis probability is always 1.

**However**, the ODE must be discretized and the **discretization error** makes the decision non-trivial.

# Hamiltonian Monte Carlo III

Finally, the ODE admits a separable stationary measure proportional to

$$\pi(x, v) \propto \exp(-H(x, v))$$

The marginal density with respect to  $x$  is therefore

$$\pi(x) = \int_{\mathbb{R}^d} \pi(x, v) dv \propto \exp(-f(x))$$

# Langevin Dynamics I

Another method for sampling is via solving the Langevin Stochastic Differential Equation which is the Newton's Second Law together with a Brownian Motion  $W$ .

$$\dot{x} = v$$

$$\dot{v} = -\gamma v - \nabla f(x) + \sqrt{2\epsilon\gamma}\dot{W}$$

where  $\dot{W}$  is the derivative of the Brownian motion, that is  $dW \sim \mathcal{N}(0, dt)$ . Under mild conditions the SDE accepts a stationary measure proportional to  $\exp\left(-\frac{1}{2}\|v\|^2 - f(x)\right)$ . The parameters  $\gamma$  (damping factor),  $\epsilon$  determine the nature of the dynamics

- ① when  $\gamma > 1$  the system is overdamped (OLD equation)
- ② when  $\gamma < 1$  the system is underdamped (ULD equation)
- ③ when  $\gamma = 1$  the system is critically damped

# Langevin Dynamics II

Of particular interest is the ULD equation

$$\begin{aligned}\dot{x} &= v \\ \dot{v} &= -2v - u\nabla f(x) + 2\sqrt{u}\dot{W} \\ u &= 1/L\end{aligned}$$

which is widely used in log-concave sampling. (for more information see [LST20, LSV18, GP14]).

# Sampling Applications

- 1 Integral Calculation (Volume Calculation etc.)
- 2 Control systems
- 3 Generative Adversarial Networks
- 4 Logistic Regression
- 5 Financial Modeling
- 6 Probabilistic Graphical Models

# ODE Solvers



# General Setting I

Our goal is to solve an ODE of the form

$$\dot{x}(t) = F(x(t), t) \quad x(0) = x_0$$

## Theorem

*If  $F$  is Lipschitz continuous in  $x$  and continuous in  $t$  then the above has a unique solution  $x(t) = \phi(t)$*

The HMC equations have  $F(x(t), v(t), t) = \begin{pmatrix} v(t) \\ -\nabla f(x(t)) \end{pmatrix}$  which is  $\max\{1, L\}$ -Lipschitz (continuous) since  $f$  is  $L$ -smooth and  $v(t)$  is 1-Lipschitz

## General Setting II

In a discrete setting the equation is solved at discrete timesteps  $t_n = t_{n-1} + \eta$  where  $\eta > 0$  is the step-size.

Let  $x_n$  denote the solution provided by the discrete solver at step  $n$  and  $\phi_n = \phi(t_n)$  be the “ideal point” at step  $n$

We define the **error**  $\epsilon_n$  to be

$$\epsilon_n = x_n - \phi_n$$

The dynamical behaviour of  $\{\epsilon_n\}_{n \geq 0}$  provides insights regarding the methods' accuracy.

# Euler Solver

The Euler Solver is the simplest one

$$t_n = t_{n-1} + \eta$$

$$x_n = x_{n-1} + \eta F(x_n, t_n)$$

It can be proven that

$$\|\epsilon_n\| \leq \frac{\eta m}{2L} (\exp(t_n - t_0) - 1) = K(t_n) \cdot \eta$$

# Runge-Kutta Methods I

The idea is to “break” every step of size  $\eta$  to smaller sub-steps and interpolate to find the next position. Each Runge-Kutta (RK) method is given by the following table (Butcher Tableau)

0				
$c_2$	$a_{21}$			
$c_3$	$a_{31}$	$a_{32}$		
$\vdots$				
$c_m$	$a_{m1}$	$\dots$	$a_{m,m-1}$	
	$b_1$	$\dots$	$b_{m-1}$	$b_m$

Table: Butcher's Tableau

where  $\sum_{j=1}^m b_j = 1$  and  $c_j = \sum_{r=1}^{j-1} a_{jr}$

# Runge-Kutta Methods II

The RK iteration proceeds in sub-steps where

$$t_n^j = t_{n-1} + c_j \eta \quad j \in [m]$$

$$k_j = F \left( \sum_{r=1}^{j-1} a_{j,r} k_r, t_n^j \right)$$

$$x_{n+1} = \sum_{j=1}^m b_j k_j$$

$$t_{n+1} = t_n + \eta$$

The global truncation error  $\|\epsilon_n\|$  is  $O(\eta^m)$ .

# Collocation Methods I

The collocation method assumes that the solution is locally approximated as

$$x(t) = \sum_{j=0}^m a_j \phi_j(t) \quad (1)$$

where  $\{\phi_j\}_{0 \leq j \leq m}$  are basis functions (e.g. polynomials). The constants  $\{a_j\}_{0 \leq j \leq m}$  are found by interpolation on the derivative of  $x$  at points given by  $t_{n+1}^j = t_n + c_j \eta$  as in the RK methods.

# Collocation Methods II

The system of equations for the interpolation is given by

$$\begin{aligned}t_{n+1}^j &= t_n + c_j \eta \\ a_{n+1}^0 &= x_n \\ \dot{x}_{n+1}^j &= F(x_{n+1}^j)(c_j - c_{j-1})\eta \\ \dot{x}_{n+1}^j &= \sum_{j=0}^m a_{n+1}^j \phi_{n+1}^j(t_{n+1}^j)\end{aligned}$$

Alternatively one solves an  $m \times m$  system of the form  $\dot{\Phi}_{n+1} a_{n+1} = \dot{X}_{n+1}$ . If the matrix of the basis derivatives is not full-rank then a solution to  $\min_{a_{n+1}} \frac{1}{2} \|\dot{\Phi}_{n+1} a_{n+1} - \dot{X}_{n+1}\|_2^2$  is sought (e.g. using SVD).

# Collocation Methods III

As choices for bases one has many choices, some of which being

- ① Polynomials  $\phi_n^j(t) = (t - t_n)^j$
- ② Lagrange polynomials  $\phi_n^j(t) = \prod_{r \neq j} \frac{t - t_r}{t_j - t_r}$
- ③ Rational functions  $\phi_n^j(t) = \frac{p_n^j(t)}{q_n^j(t)}$  with  $q_n^j(t) \neq 0$  in the ROIs.



# Boundary Oracles

# Boundary Conditions I

In HMC the domain of  $(x, v)$  is  $K \times \mathbb{R}^d \subseteq \mathbb{R}^d \times \mathbb{R}^d$ . Where  $K \neq \mathbb{R}^d$  one has to account for **boundary conditions** for the position  $x$ .

There are three main types of boundary conditions

- 1 Neumann Conditions (Boundary Reflections) where  $\frac{\partial x}{\partial n} = 0$
- 2 Dirichlet Conditions  $x = g$
- 3 Robin (mixed) Conditions  $a \frac{\partial x}{\partial n} + g = 0$

where the domain of  $a, g$  is the boundary  $\partial K$ .

It has been proven [PP14] that HMC admits boundary conditions equivalent to the **Neumann Conditions**.

# Boundary Conditions II




Figure: Baseline and Reflective HMC. Taken from [AD15].

# The Reflection Operator I

A point  $x$  reflects at the boundary point  $\tilde{x}$  with normal  $n$ .

We define the reflection operator  $\text{refl}$  such that

$$\text{refl}(x) = -2(a^T n)n + a + \tilde{x}$$

where  $a = \tilde{x} - x$  is the ray between the initial and the boundary points. Note that in general  $\text{refl}(x)$  may not lie in  $K$ . We compose the reflection operator  $k$  times such that  $\text{refl}^k(x) = \text{refl} \circ \dots \circ \text{refl}(x) \in K$ . In our setting we assume that at each step the proposal point cannot reflect more than  $\ell \in \mathbb{N}^*$  times.

# The Reflection Operator II




Figure: Reflection Illustration.  $x_1'$  is the reflection of  $x_1$  about  $x_0$  with normal  $n$ .

Source: <https://mathworld.wolfram.com/Reflection.html>

# Computing Intersections with $\partial K$ I

Of particular interest is the computation of the intersection of an (implicit) curve between a point  $x$  inside the convex body  $K$  and a proposal  $\tilde{x} \notin K$ .

**Case 1.** The curve is a *line segment* and  $K$  is a convex polytope.

We parametrize the line segment between  $x$  and  $\tilde{x}$  with

$\gamma(t) = tx + (1 - t)\tilde{x}$  where  $t \in [0, 1]$ . We seek

$t_u = \sup\{t \in [0, 1] | \gamma(t) \in \partial K\}$  and  $u = \gamma(t_u)$  as the solution to the boundary intersection problem.

We use the Cyrus-Beck [CB78] algorithm

## Computing Intersections with $\partial K$ II

Let  $z \in \partial K$  be known and let  $n$  represent the normal vector at  $z$ . We compute the quantity

$$n^T(\gamma(t) - z) \begin{cases} = 0 & \gamma(t) \in \partial K \\ < 0 & \gamma(t) \notin K \\ > 0 & \gamma(t) \in K \setminus \partial K \end{cases}$$

Solving the equation  $n^T(\gamma(t) - z)$  for  $t$  we get

$$t = \frac{n^T(z - x)}{n^T(\tilde{x} - x)}$$

We compute the above for all the  $N$  normals of the polytope and keep the maximum value that lies in  $[0, 1]$ . The min value can also be kept in case we want the other intersection point as well. Complexity is  $O(Nd)$

## Computing Intersections with $\partial K$ III

**Case 2.** The curve has the form  $\gamma(t) = \sum_{i=1}^m a_i \phi_i(t)$ ,  $\{\phi_j\}_{j \in [m]}$  are basis functions, and  $K$  is a convex polytope.

We use the same procedure as above, however now we cannot solve directly for  $t$ . We, for example, can use the Newton-Raphson root finder to solve the transcendental equation.

$$t^{(r+1)} = t^{(r)} - \frac{\sum_{j \in [m]} (n^T a_j) \phi_j(t^{(r)}) - n^T z}{\sum_{j \in [m]} (n^T a_j) \dot{\phi}_j(t^{(r)})} \quad r \in [R]$$

Complexity is  $O(NdR)$  where  $R$  is the maximum number of iterations the NR solver must be called to find a root.

**Problems.** Convergence, Well-posedness (denominator getting too small)



## Computing Intersections with $\partial K$ IV

**Case 3.** The convex body  $K$  has the form  $K = \{x \in \mathbb{R}^d | g(x) \leq 0\}$  where  $g(x) = \max_{1 \leq i \leq M} g_i(x)$  where  $g_1, \dots, g_M$  are twice-differentiable convex functions that are  $\mu$ -strongly-convex.

**Examples.**  $L_2$  Balls, Spectrahedra etc.

**Idea.** Linearize the convex body around  $x + h$

$$0 \geq g_i(x + h) \geq g_i(x) + \langle \nabla g_i(x), h \rangle + \frac{\mu \|h\|^2}{2}$$

The linearized convex polytope  $P(x)$  around  $x$  is

$$J(x)h \leq b$$

where  $J(x)$  is the Jacobian matrix around  $x$  with entries  $J_{ij}(x) = \frac{\partial g_i(x)}{\partial x_j}$  and  $b$  has entries  $b_i = -g_i(x)$ .

# Computing Intersections with $\partial K$

The linear approximation error is  $O(\|h\|^2)$ . A high-level algorithm proceeds as follows.

We are given a curve  $\gamma(t)$  and a starting point  $x_0 = \gamma(0)$ , an accuracy  $\epsilon > 0$ , and a step counter  $i$  initialized at 0.

- 1 Find  $P(x_i)$  around  $x_i$  and the intersection point of  $\gamma(t)$  with  $P(x_i)$  (see Case 1, Case 2). Let that point be  $x_{i+1} = \gamma(t_{i+1})$
- 2 Calculate  $g(x_{i+1}) = \max_{1 \leq j \leq M} g_j(x_{i+1})$ . If  $|g(x_{i+1})| \leq \epsilon$ , output  $x_{i+1}, t_{i+1}$ , else repeat.

# Next Talk(s)

Next talk(s) will be occupied with

- 1 Algorithmic Issues for the sampling problem (mixing time, bounds etc.).
- 2 Theoretical contributions to the problem.
- 3 Implementation details.

# Thank you!

# References I



Hadi Mohasel Afshar and Justin Domke.

Reflection, refraction, and hamiltonian monte carlo.

In *Advances in neural information processing systems*, pages 3007–3015, 2015.



Mike Cyrus and Jay Beck.

Generalized two-and three-dimensional clipping.

*Computers & Graphics*, 3(1):23–28, 1978.



Elena Gryazina and Boris Polyak.

Random sampling: Billiard walk algorithm.

*European Journal of Operational Research*, 238(2):497–504, 2014.

# References II



Yin Tat Lee, Ruoqi Shen, and Kevin Tian.

Logsmooth gradient concentration and tighter runtimes for metropolized hamiltonian monte carlo.

*arXiv preprint arXiv:2002.04121*, 2020.



Yin Tat Lee, Zhao Song, and Santosh S Vempala.

Algorithmic theory of odes and sampling from well-conditioned log-concave densities.

*arXiv preprint arXiv:1812.06243*, 2018.



Ari Pakman and Liam Paninski.

Exact hamiltonian monte carlo for truncated multivariate gaussians.

*Journal of Computational and Graphical Statistics*, 23(2):518–542, 2014.