

Log-concave Sampling (Part 1)


Marios Papachristou

GeomScale, NTUA

June 3, 2020



GeomScale



Mentors: A. Chalkis (NKUA), V. Fisikopoulos (NKUA), E. Tsigaridas (INRIA)

Homepage: https://github.com/GeomScale/volume_approximation

About today's talk / tutorial

Today's talk will concentrate on

Sampling from high-dimensional log-concave densities

- 1 Introduction to log-concave sampling.
- 2 ODE Solvers.
- 3 Boundary Oracles.

Google Summer of Code 2020

The current GSoC project aims to provide implementations (and theoretical insights) to log-concave sampling problems for the GeomScale project.

Milestones

① Milestone I (ODE Solvers)

- Implement ODE solvers (Euler, Runge-Kutta, Collocation, etc.)
- Efficiently address boundary oracles

② Milestone II (Samplers)

- Implement samplers (HMC, Langevin etc.).
- Provide theoretical guarantees on truncated settings.

③ Milestone III (R bindings)

- Port C++ functionality of

Today's talk will mostly concentrate on **Milestone I**.

Our project involves taking samples from distributions with probability density functions of the form

$$\pi(x) \propto \exp(-f(x)) \quad x \in K$$

where K is either: (a) \mathbb{R}^d , or (b) a convex body, and f is a convex function that is L -smooth and m -strongly convex.

Convex Functions I

A domain K is convex iff (if and only if) for all $x, y \in K$ it holds that for all $t \in [0, 1]$

$$tx + (1 - t)y \in K$$

The domain K is a convex body iff it is convex, closed and bounded.

A function $f : K \rightarrow \mathbb{R}$ is convex iff for all $x, y \in K$ we have that for all $t \in [0, 1]$

$$f(tx + (1 - t)y) \leq tf(x) + (1 - t)f(y)$$

Convex functions have some very nice properties, and their use is widespread in optimization.

Convex Functions II

If the function is twice differentiable with gradient ∇f and Hessian matrix $\nabla^2 f$ then

- We say that f is L -smooth iff $\|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|$ or $\nabla^2 f(x) \preceq L \cdot I_d$.
- We say that f is m -strongly convex iff

$$f(y) \geq f(x) + \langle \nabla f(x), y - x \rangle + \frac{m}{2} \|x - y\|^2$$

or $\nabla^2 f(x) \succeq m \cdot I_d$.

- We define the **condition number** of f to be the ratio of max/min eigenvalues of the Hessian, that is $\kappa = L/m$.

Log-concave Sampling I

Our goal is sampling from $\pi(x) \propto \exp(-f(x))$.

Directly sampling from $\pi(x)$ is very difficult since one has to account for the normalization constant $\int_K \exp(-f(x)) dx$ which is in general **intractable**.

Idea. The distribution $\pi(x)$ can be thought as the stationary measure of a Markov Chain that is $\pi(x) = \lim_{k \rightarrow \infty} \pi_k(x)$.

Log-concave Sampling II

One of the first algorithms to do it is the Metropolis-Hastings Algorithm. The general idea of Metropolis Hastings is

- 1 Assume that you are at a state x
- 2 Perform a transition to a new nearby state y and make a proposal
- 3 Accept the proposal to move to y with probability (Metropolis Filter)

$$\min \left\{ 1, \frac{a(x, y)\pi(y)}{a(y, x)\pi(x)} \right\}$$

where a is a transition probability function.

It can be shown analytically that the above process converges to a stationary distribution $\pi(x)$.

Intuition when $a(x, y) = a(y, x)$: The sampler has incentive to move towards higher-density areas (but lower density areas are also allowed)

Algorithmic Challenges

The main Algorithmic Challenges for (log-concave) sampling are in general

- 1 The **Mixing Time** of the Markov Chain, that is how fast ($\#$ iterations) a Markov Chain with transition operator \mathcal{T} starting from an initial distribution π_0 reaches π within Total Variation Distance of at most $\delta > 0$ (more next time)

$$t_{\text{mix}}(\delta) = \inf \{k \geq 0 \mid \|\mathcal{T}^k(\pi_0) - \pi\|_{TV} \leq \delta\}$$

where $\|P - Q\|_{TV} = \sup_{A \in \mathcal{F}} |P(A) - Q(A)|$ and \mathcal{F} is a σ -algebra on the state space K .

- 2 The **Cost-Per Iteration** that is

$$\left(\begin{array}{c} \text{Cost-per} \\ \text{iteration} \end{array} \right) = \left(\begin{array}{c} \text{Cost-per} \\ \text{ODE step} \end{array} \right) + \left(\begin{array}{c} \text{Cost-per boundary} \\ \text{oracle (if truncated)} \end{array} \right)$$

Sampling in a continuous setting

Hamiltonian Monte Carlo I

The state-space is continuous and the samples can be proposed via solving Hamilton's equations for a particle with position x and velocity v under a conservative potential $f(x)$ that applies a force $-\nabla f(x)$.

The Hamiltonian of the particle is defined as

$$H(x, v) = \frac{1}{2} \|v\|^2 + f(x)$$

and Hamilton's equations simulate the particle's behaviour

$$\begin{aligned}\dot{x} &= \frac{\partial H}{\partial v} = v \\ \dot{v} &= -\frac{\partial H}{\partial x} = -\nabla f(x)\end{aligned}$$

Hamiltonian Monte Carlo II

We start by choosing a direction $v \sim \mathcal{N}(0, I_d)$ and simulate one/many steps of the ODE arriving at a proposal (\tilde{x}, \tilde{v}) .

The Metropolis Filter in this case for a proposal (\tilde{x}, \tilde{v}) given a state (x, v) is $\min\{1, \exp(H(\tilde{x}, \tilde{v}) - H(x, v))\}$.

Ideally note that $\dot{H} = \langle \nabla_{x,v} H, (\dot{x}, \dot{v}) \rangle = \langle (v, \nabla f(x)), (-\nabla f(x), v) \rangle = 0$ and hence the Metropolis probability is always 1.

However, the ODE must be discretized and the **discretization error** makes the decision non-trivial.

Hamiltonian Monte Carlo III

Finally, the ODE admits a separable stationary measure proportional to

$$\pi(x, v) \propto \exp(-H(x, v))$$

The marginal density with respect to x is therefore

$$\pi(x) = \int_{\mathbb{R}^d} \pi(x, v) dv \propto \exp(-f(x))$$

Langevin Dynamics I

Another method for sampling is via solving the Langevin Stochastic Differential Equation which is the Newton's Second Law together with a Brownian Motion W .

$$\begin{aligned}\dot{x} &= v \\ \dot{v} &= -\gamma v - \nabla f(x) + \sqrt{2\epsilon\gamma} \dot{W}\end{aligned}$$

where \dot{W} is the derivative of the Brownian motion, that is $dW \sim \mathcal{N}(0, dt)$. Under mild conditions the SDE accepts a stationary measure proportional to $\exp(-\frac{1}{2}\|v\|^2 - f(x))$. The parameters γ (damping factor), ϵ determine the nature of the dynamics

- ① when $\gamma > 1$ the system is overdamped (OLD equation)
- ② when $\gamma < 1$ the system is underdamped (ULD equation)
- ③ when $\gamma = 1$ the system is critically damped

Langevin Dynamics II

Of particular interest is the ULD equation

$$\begin{aligned}\dot{x} &= v \\ \dot{v} &= -2v - u \nabla f(x) + 2\sqrt{u} \dot{W} \\ u &= 1/L\end{aligned}$$

which is widely used in log-concave sampling. (for more information see [LST20, LSV18, GP14]).

Sampling Applications

- 1 Integral Calculation (Volume Calculation etc.)
- 2 Control systems
- 3 Generative Adversarial Networks
- 4 Logistic Regression
- 5 Financial Modeling
- 6 Probabilistic Graphical Models

ODE Solvers

General Setting I

Our goal is to solve an ODE of the form

$$\dot{x}(t) = F(x(t), t) \quad x(0) = x_0$$

Theorem

If F is Lipschitz continuous in x and continuous in t then the above has a unique solution $x(t) = \phi(t)$

The HMC equations have $F(x(t), v(t), t) = \begin{pmatrix} v(t) \\ -\nabla f(x(t)) \end{pmatrix}$ which is Lipschitz (continuous) since f is L -smooth and $v(t)$ is 1-Lipschitz

General Setting II

In a discrete setting the equation is solved at discrete timesteps $t_n = t_{n-1} + \eta$ where $\eta > 0$ is the step-size.

Let x_n denote the solution provided by the discrete solver at step n and $\phi_n = \phi(t_n)$ be the “ideal point” at step n

We define the **error** ϵ_n to be

$$\epsilon_n = x_n - \phi_n$$

The dynamical behaviour of $\{\epsilon_n\}_{n \geq 0}$ provides insights regarding the methods' accuracy.

Euler Solver

The Euler Solver is the simplest one

$$\begin{aligned}t_n &= t_{n-1} + \eta \\x_n &= x_{n-1} + \eta F(x_{n-1})\end{aligned}$$

It can be proven that

$$\|\epsilon_n\| \leq \frac{\eta m}{2L} (\exp(t_n - t_0) - 1) = K(t_n) \cdot \eta$$

Hence the error of the Euler Solver is $O(\eta)$.

Runge-Kutta Methods I

The idea is to “break” every step of size η to smaller sub-steps and interpolate to find the next position. Each Runge-Kutta (RK) method is given by the following table (Butcher Tableau)

0				
c_2	a_{21}			
c_3	a_{31}	a_{32}		
\vdots				
c_m	a_{m1}	\dots	$a_{m,m-1}$	
	b_1	\dots	b_{m-1}	b_m

Table: Butcher's Tableau

where $\sum_{j=1}^m b_j = 1$ and $c_j = \sum_{r=1}^{j-1} a_{jr}$

Runge-Kutta Methods II

The RK iteration proceeds in sub-steps where

$$\begin{aligned}t_n^j &= t_{n-1} + c_j \eta & j \in [m] \\k_j &= F \left(\sum_{r=1}^{j-1} a_{j,r} k_r, t_n^j \right) \\x_{n+1} &= \sum_{j=1}^m b_j k_j \\t_{n+1} &= t_n + \eta\end{aligned}$$

The global truncation error $\|\epsilon_n\|$ is $O(\eta^m)$.

Collocation Methods I

The collocation method assumes that the solution is locally approximated as

$$p(t) = \sum_{j=0}^m a_j \phi_j(t) \quad (1)$$

where $\{\phi_j\}_{0 \leq j \leq m}$ are basis functions (e.g. polynomials). The constants $\{a_j\}_{0 \leq j \leq m}$ are found by interpolation on the derivative of x at points given by $t_{n+1}^j = t_n + c_j \eta$ as in the RK methods.

As choices for bases one has many choices, some of which being

- ① Polynomials $\phi_n^j(t) = (t - t_n)^j$
- ② Lagrange polynomials $\phi_n^j(t) = \prod_{r \neq j} \frac{t - t_r}{t_j - t_r}$
- ③ Rational functions $\phi_n^j(t) = \frac{p_n^j(t)}{q_n^j(t)}$ with $q_n^j(t) \neq 0$ in the ROIs.

Collocation Methods II

The system of equations for the interpolation is given by

$$\begin{aligned}t_{n+1}^j &= t_n + c_j \eta \\p_{n+1}(t_{n+1}^0) &= x_n \\ \dot{p}_{n+1}(t_{n+1}^j) &= F(p_{n+1}(t_{n+1}^j)) \quad j \in [m]^*\end{aligned}$$

Alternatively if F is a linear mapping one solves an $m \times m$ system of the form $\dot{\Phi}_{n+1} a_{n+1} = \dot{X}_{n+1}$. If the matrix of the basis derivatives is not full-rank then a solution to $\min_{a_{n+1}} \frac{1}{2} \|\dot{\Phi}_{n+1} a_{n+1} - \dot{X}_{n+1}\|_2^2$ is sought (e.g. using SVD).

If F is non-linear one can in general use iterative methods (NR) to compute the coefficients $\{a_{n+1}^j\}_{n \geq 0, j \in [m]}$

Leapfrog Integrator (2nd order)

The Leapfrog integrator is used to solve the equation $\ddot{x} = F(x, t)$

The method proceeds as follows

$$v_{i+1/2} = v_i + \frac{\eta}{2} F(x_i)$$

$$x_{i+1} = x_i + \eta v_{i+1/2}$$

$$v_{i+1} = v_{i+1/2} + \frac{\eta}{2} F(x_{i+1})$$

Examples of interest. Particle dynamics (HMC, Langevin)

Boundary Oracles

Boundary Conditions I

In HMC the domain of (x, v) is $K \times \mathbb{R}^d \subseteq \mathbb{R}^d \times \mathbb{R}^d$. Where $K \neq \mathbb{R}^d$ one has to account for **boundary conditions** for the position x .

There are three main types of boundary conditions

- 1 Neumann Conditions (Boundary Reflections) where $\frac{\partial x}{\partial n} = 0$
- 2 Dirichlet Conditions $x = g$
- 3 Robin (mixed) Conditions $a \frac{\partial x}{\partial n} + g = 0$

where the domain of a, g is the boundary ∂K .

It has been proven [PP14] that HMC admits boundary conditions equivalent to the **Neumann Conditions**.

Boundary Conditions II




Figure: Baseline and Reflective HMC. Taken from [AD15].

The Reflection Operator I

A point x reflects at the boundary point \tilde{x} with normal n .

We define the reflection operator refl such that

$$\text{refl}(x) = -2(a^T n)n + a + \tilde{x}$$

where $a = \tilde{x} - x$ is the ray between the initial and the boundary points. Note that in general $\text{refl}(x)$ may not lie in K . We compose the reflection operator k times such that $\text{refl}^k(x) = \text{refl} \circ \cdots \circ \text{refl}(x) \in K$. In our setting we assume that at each step the proposal point cannot reflect more than $\ell \in \mathbb{N}^*$ times.

The Reflection Operator II




Figure: Reflection Illustration. x_1' is the reflection of x_1 about x_0 with normal n . Source: <https://mathworld.wolfram.com/Reflection.html>

Computing Intersections with ∂K I

Of particular interest is the computation of the intersection of an (implicit) curve between a point x inside the convex body K and a proposal $\tilde{x} \notin K$.

Case 1. The curve is a *line segment* and K is a convex polytope.

We parametrize the line segment between x and \tilde{x} with $\gamma(t) = tx + (1 - t)\tilde{x}$ where $t \in [0, 1]$. We seek $t_u = \sup\{t \in [0, 1] | \gamma(t) \in \partial K\}$ and $u = \gamma(t_u)$ as the solution to the boundary intersection problem.

We use the Cyrus-Beck [CB78] algorithm

Computing Intersections with ∂K II

Let $z \in \partial K$ be known and let n represent the normal vector at z . We compute the quantity

$$n^T(\gamma(t) - z) \begin{cases} = 0 & \gamma(t) \in \partial K \\ < 0 & \gamma(t) \notin K \\ > 0 & \gamma(t) \in K \setminus \partial K \end{cases}$$

Solving the equation $n^T(\gamma(t) - z)$ for t we get

$$t = \frac{n^T(z - x)}{n^T(\tilde{x} - x)}$$

We compute the above for all the N normals of the polytope and keep the maximum value that lies in $[0, 1]$. The min value can also be kept in case we want the other intersection point as well. Complexity is $O(Nd)$

Computing Intersections with ∂K III

H-polytope. The polytope is given by the form $Ax \leq b$ where A consists of N row vectors $a_1, \dots, a_N \in \mathbb{R}^d$ and $b = (b_1, \dots, b_N)^T$. On each facet we use a_i as normal vector and if $b_i = 0$ then we use $\vec{0}$ as a point on the facet. If $b_i \neq 0$, there exists at least one index r such that $a_r \neq 0$ (otherwise the problem is trivial) we use the point $u = (0, 0, \dots, b_i/a_r, \dots, 0)^T$ which lies on the hyperplane, that is $a_i^T u = b_i$. Worst-case complexity is $O(Nd)$.

Computing Intersections with ∂K IV

V-polytope. The polytope is given by its convex hull V which contains M points $v_1, \dots, v_M \in \mathbb{R}^d$. The point $\gamma(t) = tx + (1 - t)\tilde{x}$ is on the boundary for some $t_0 \in [0, 1]$ (given that $x \in K$) if t_0 is the maximum value of $t \in [0, 1]$ such that there exist $\lambda_1, \dots, \lambda_M \geq 0$ with $\sum_{i=1}^M \lambda_i = 1$ and $\gamma(t_0) = \sum_{i=1}^M \lambda_i v_i$, which translates to the following LP problem which has $O(Md)$ constraints

$$\begin{aligned} & \text{maximize} && t \\ & \text{subject to} && 0 \leq t \leq 1 \\ & && \lambda_i \geq 0 && i \in [M] \\ & && \sum_{i=1}^M \lambda_i = 1 \\ & && tx + (1 - t)\tilde{x} - \sum_{i=1}^M \lambda_i v_i = 0 \end{aligned}$$

Solvable via `lp_solve` (functionality already exists)

Computing Intersections with $\partial K \vee$

Case 2. The curve has the form $\gamma(t) = \sum_{i=1}^m a_i \phi_i(t)$, $\{\phi_j\}_{j \in [m]}$ are basis functions, and K is a convex polytope.

H-polytope. We use the same procedure as above, however now we cannot solve directly for t . We, for example, can use the Newton-Raphson root finder to solve the transcendental equation.

$$t^{(r+1)} = t^{(r)} - \frac{\sum_{j \in [m]} (n^T a_j) \phi_j(t^{(r)}) - n^T z}{\sum_{j \in [m]} (n^T a_j) \dot{\phi}_j(t^{(r)})} \quad r \in [R]$$

Complexity is $O(NdR)$ where R is the maximum number of iterations the NR solver must be called to find a root.

Problems. Convergence, Well-posedness (denominator getting too small)

Computing Intersections with ∂K VI

V-polytope. The problem of Case I is a general optimization problem

$$\begin{array}{ll}\text{maximize} & t \\ \text{subject to} & t \geq 0 \\ & \lambda_i \geq 0 \quad i \in [M] \\ & \sum_{i=1}^M \lambda_i = 1 \\ & \sum_{i=1}^M a_j \phi_j(t) - \sum_{i=1}^M \lambda_i v_i = 0\end{array}$$

Can be addressed using the KKT conditions.

Computing Intersections with ∂K VII

Case 3. The convex body K has the form $K = \{x \in \mathbb{R}^d \mid g(x) \leq 0\}$ where $g(x) = \max_{1 \leq i \leq M} g_i(x)$ where g_1, \dots, g_M are twice-differentiable convex functions that are μ -strongly-convex.

Examples. L_2 Balls, Spectrahedra etc.

Idea. Linearize the convex body around $x + h$

$$0 \geq g_i(x + h) \geq g_i(x) + \langle \nabla g_i(x), h \rangle + \frac{\mu \|h\|^2}{2}$$

The linearized convex polytope $P(x)$ around x is

$$J(x)h \leq b$$

where $J(x)$ is the Jacobian matrix around x with entries $J_{ij}(x) = \frac{\partial g_i(x)}{\partial x_j}$ and b has entries $b_i = -g_i(x)$.

Computing Intersections with ∂K VIII

The linear approximation error is at most $\frac{\mu}{2}\|h\|^2$. A high-level algorithm (Local-search-based) proceeds as follows.

We are given a curve $\gamma(t)$ and a starting point $x_0 = \gamma(0)$, an accuracy $\epsilon > 0$, and a step counter i initialized at 0.

- 1 Find $P(x_i)$ around x_i and the intersection point of $\gamma(t)$ with $P(x_i)$ (see Case 1, Case 2). Let that point be $x_{i+1} = \gamma(t_{i+1})$
- 2 Calculate $g(x_{i+1}) = \max_{1 \leq j \leq M} g_j(x_{i+1})$. If $|g(x_{i+1})| \leq \epsilon$, output x_{i+1}, t_{i+1} , else repeat.

Next Talk(s)

Next talk(s) will be occupied with

- 1 Algorithmic Issues for the sampling problem (mixing time, bounds etc.).
- 2 Theoretical contributions to the problem.
- 3 Implementation details.

Thank you!

References I



Hadi Mohasel Afshar and Justin Domke.

Reflection, refraction, and hamiltonian monte carlo.

In Advances in neural information processing systems, pages 3007–3015, 2015.



Mike Cyrus and Jay Beck.

Generalized two-and three-dimensional clipping.

Computers & Graphics, 3(1):23–28, 1978.



Elena Gryazina and Boris Polyak.

Random sampling: Billiard walk algorithm.

European Journal of Operational Research, 238(2):497–504, 2014.



Yin Tat Lee, Ruoqi Shen, and Kevin Tian.

Logsmooth gradient concentration and tighter runtimes for metropolized hamiltonian monte carlo.

arXiv preprint arXiv:2002.04121, 2020.

References II



Yin Tat Lee, Zhao Song, and Santosh S Vempala.

Algorithmic theory of odes and sampling from well-conditioned log-concave densities.

arXiv preprint arXiv:1812.06243, 2018.



Ari Pakman and Liam Paninski.

Exact hamiltonian monte carlo for truncated multivariate gaussians.

Journal of Computational and Graphical Statistics, 23(2):518–542, 2014.